The complex Grothendieck constant

New bounds, machine-checked in Lean

Alin Galatan, with several AI research systems

https://github.com/agalatan/grothendieck-lean

October 2026

$$1.401697117301 \;<\; K_G^{\mathbb{C}} \;<\; 1.404598149141$$

Both bounds are theorems of a Lean 4/Mathlib development, resting on explicitly stated analytic and computer-assisted inputs (Section 3).

Only the sharp upper target is active; its fourteen finite-band Gaussian estimates and energy-table accuracy remain to be proved.

The source-node rounding ledger verifies its exact rational arithmetic and transfers a squared-mass cap to a norm cap. Its application to the actual Gaussian coefficient table still requires the input estimates and the operation-error model; it does not complete the upper endpoint.

Abstract

Grothendieck's inequality bounds how much a semidefinite (vector) relaxation of a bilinear optimisation problem can exceed the true optimum. The best constant in the complex case, $K_G^{\mathbb{C}}$, has been known since 1987 to lie roughly in $[1.338, 1.405]$; its first two decimals were settled only in September 2026. We improve both bounds, to an interval of width 0.0029, and verify the arguments in Lean. This note explains the problem, the two constructions, exactly what the computer checks, the new inequalities and obstructions found along the way, and why the remaining gap is hard: the two sides are dual optimisation problems, and the obstruction on both is the nonlinear (cubic and higher) structure of bounded functions of Gaussian vectors.

1 The question

Let $A = (a_{ij})$ be an $m \times n$ complex matrix. Compare

$$C(A) = \sup_{|\varepsilon_i|=|\delta_j|=1} \Big|\sum_{i,j} a_{ij}\,\overline{\varepsilon_i}\,\delta_j\Big|, \qquad V(A) = \sup_{\|x_i\|=\|y_j\|=1} \Big|\sum_{i,j} a_{ij}\,\langle x_i, y_j\rangle\Big|,$$

the first over unimodular scalars, the second over unit vectors in any complex Hilbert space. Clearly $C(A) \le V(A)$. Grothendieck proved in 1953 that $V(A) \le K\,C(A)$ for a constant K independent of m, n and the dimension. The smallest such K is the *complex Grothendieck constant* $K_G^{\mathbb{C}}$.

Why it matters.

- *Optimisation.* $V(A)$ is a semidefinite program; $C(A)$ is a hard combinatorial problem. $K_G^{\mathbb{C}}$ is the best approximation ratio of the semidefinite relaxation, and rounding algorithms for it are exactly upper-bound proofs.

- *Functional analysis.* K_G is the norm of the identity between the injective and projective tensor norms on $\ell_\infty \otimes \ell_\infty$, and controls absolutely summing operators (Grothendieck's original setting).
- *Quantum information.* The real constant $K_G^{\mathbb{R}}$ is the largest ratio between quantum and classical values of two-player XOR games (Tsirelson); the finite-order constants $K_G(d)$ govern the nonlocality of qubit states.

History.

year	who	lower bound	upper bound
1953	Grothendieck		$K < \infty$
1984	Davie	1.33807	
Sep 2026	Guo, Fang, Lu	1.35584631827	
Sep 2026	Heilman, Jones, Malavolta	1.4	archived
Oct 2026	**this work**	**1.401697117301**	**1.404598149141**

Haagerup's integral $1/\gamma^\dagger = 1.404575934\ldots$, $\gamma^\dagger = \int_0^{\pi/2} \cos^2 t\,(1+\sin^2 t)^{-1/2} dt$, has been proposed as the exact value. It lies inside our interval, $2.22 \cdot 10^{-5}$ below the upper bound. For the real constant the corresponding recent results are $1.773 \le K_G^{\mathbb{R}} \le 1.7799$ (Heilman–Jones–Malavolta, October 2026).

2 Two dual constructions

Both bounds use the same Gaussian picture. Let $Z = (z_1, \ldots, z_n)$ be a standard complex Gaussian vector and decompose $L^2(\mathbb{C}^n, \gamma)$ into the *Wick chaoses* $H_{p,q}$ (polynomials of bidegree (p, q) in $z, \bar{z}$, orthogonalised). Under the rotation $z \mapsto e^{i\theta} z$ every relevant object has *charge one*, so only the bidegrees $(1,0), (2,1), (3,2), \ldots$ (degrees $1, 3, 5, \ldots$) appear. The standard complex Gaussian measure, its rotation invariance and radial-moment reduction form a shared Lean foundation. The upper rounding proof depends on this foundation independently of the lower row certificates.

2.1 Upper bounds: rounding schemes

The formal quadrature development includes exact polynomial identities for each original Gauss order from 1 to 64. Rational polynomial certificates, interpreted through LeanCert, prove orthogonality and the weight identity; these give exact integration for polynomials of degree below $2n$ at order n. The certificates use kernel reduction and are part of the main Lean build. The selected numerical nodes and the analytic remainder require separate proofs; this algebraic result alone does not finish the upper endpoint. Given unit vectors, project a Gaussian onto them and round: $\varepsilon_i = f(\langle G, x_i\rangle)$ with $|f| \le 1$. The correlation of two rounded values depends only on $t = \langle x_i, y_j\rangle$,

$$\mathbb{E}[\overline{f(\cdot)}\, g(\cdot)] = H(t) = \sum_{k \ge 0} c_k\, t|t|^{2k}.$$

If $H(t) = c\,t + r(t)$ with $\|r\|_A = \sum_k |c_k - c\,\delta_{k0}| < c$, a tensor-power argument (Krivine's trick) gives $K_G^{\mathbb{C}} \le 1/(c - \|r\|_A)$. The shared Gaussian phase identity is proved from the Hermite basis and Mehler formula. Our sharp rounding is a smoothed two-coordinate phase rounding, $F_0 = \mathrm{phase}(z + \sqrt{1+a^2}\, v(|z|^2)\, w)$ with a twelve-term Laguerre profile v, multiplied by the cubic Taylor polynomial of a damped polynomial phase φ: $T = F_0 P_3(\varphi)$, $P_3(x) = 1 + ix - x^2/2 - ix^3/6$. The identity $|P_3(x)|^2 = 1 - x^4(3-x^2)/36$ and the proved bound $\varphi^2 \le 99/100$ (hence $|\varphi| < 1$) give $|T| \le 1$, and a disk-valued rounding is as good as a unimodular one. The phase certificate checks

119,077 leaves in 512 cubes through six native range facts, with an analytic far field. Its kernel is $H(t) = \sum E[n,m]\, p(t)^n(-\tau(t))^m$ with an odd polynomial p of degree 41 and τ of degree 15 with zero linear coefficient (any odd preprocessors of Wiener norm ≤ 1 are realised by tensor powers). Compared with $c\,t$, $c = 7119483103461/10^{13}$, the residual is at most $\delta = 455212621809/(5\cdot 10^{17}) \approx 9.10425\cdot 10^{-7}$:

$$K_G^{\mathbb{C}} \le \frac{1}{c-\delta} = \frac{500000000000000000}{355973699960428191} < 1.404598148840.$$

The whole difficulty is to keep *every* coefficient c_k, $k \ge 1$, small at once: improving the cubic coefficient c_1 tends to spoil $c_2, c_3, \dots$, and the infinite tail has to be bounded too.

The conditional-node table has 21,577 rows of binary64 conditional Hermite approximations. Row i uses the degree-531 canonical pairs $0 \le c \le 266$, $0 \le d \le 265$, $c \le d+1$, in $(c+d, c)$ order, retaining $c + d \le \text{cutoff}_i$. A full row has 35,777 entries; cutoffs 1 and 2 retain three and five. This certificate selects every node in a panel where the outward 128-bit profile interval Horner screen contains zero at some node. The 413 unresolved nodes occupy 381 panels; panel closure gives exactly the 417 ordinals 3872–4036, 10910–11037 and 17321–17444, independently of cutoff. The rule and its exact inputs are committed in `data/q0_quadrature/conditional_source/exception-selection.json`. Lean checks scalar magnitude at most one, squared row norm at most 121/100, and source ordinal coverage. The source-error assembly uses these rows through their source-address and cutoff binding; production basis/sign normalization and actual approximation estimates remain separate obligations.

The exceptional source panels have Gauss orders one through four. Exact normalization of their selected roots removes the panel-weight-sum input, using the same physical weights and one Gaussian density factor. The selected coefficient error is at most $2.53\cdot 10^{-10}$ given a 10^{-12} enclosure for the actual regular-node coefficients. Exact reconstruction of the full conditional rectangle multiplies this allowance by 10/7 and retains the production phase. Regular evaluator bounds, floating operations, contraction, omitted-domain estimates, tails and the final energy-table fit are still separate requirements for the upper endpoint.

For selected regular nodes near a profile zero, the whole Gaussian slice remainder is projected by Bessel before aggregating coefficients. Its vector radius remains $5\rho^2$, without a mode-count factor. The two-coordinate parameter error has radius at most $|\rho - r|$; a separate coupled saved-model radius ε gives $5\rho^2 + |\rho - r| + \varepsilon$. The hybrid source budget combines those selected nodes with explicit squared-error budgets on all other regular nodes, using the same physical Gaussian weights. Actual parameter and saved-vector certificates remain required.

For an exceptional source row, the model-to-payload discrepancy over all selected modes reduces exactly to five rational coordinates. The fitted parameter is $x_{10} - x_{01}$, with squared residual $(1-x_{00})^2 + (x_{10}+x_{01})^2/2 + x_{02}^2 + x_{11}^2$. Any other parameter r adds $(r - (x_{10}-x_{01}))^2/2$. The decoded source ordinal, cutoff and rank lookup are retained. Physical Gaussian parameter enclosures remain required; only the stored approximation and explicit model vanish outside the five-coordinate support. Separate kernel checks on the decoded rational rows give model residual norm below 10^{-31} and fitted-parameter magnitude below 10^{-15} at every exceptional source row. Three centered interval checks on the original exceptional panel blocks bound the actual scaled profile by $B = 1/(2\cdot 10^{14})$. Source geometry gives $\sqrt{s} \ge 3$, hence physical ratio at most $B/3$ and ratio-to-fit error at most $B/3 + 10^{-15}$. The complete selected-vector radius is then $5(B/3)^2 + B/3 + 10^{-15} + 10^{-31} < 3\cdot 10^{-15}$. The new consumer has no numerical enclosure premise. The original discrete Gaussian weights then give exceptional squared error at most $(3\cdot 10^{-15})^2 \cdot 2\cdot 10^{-21} \le (2\cdot 10^{-25})^2$, for every finite subset of the selected modes. The full-source adapter adds $(2\cdot 10^{-25})^2$ to an explicit regular squared-error budget; continuous root-gap omission is separate. The regular nodes, quadrature and remaining upper-bound obligations are separate.

For each nonnegative squared radial node, the finite auxiliary-integral response has the actual affine argument of the base F_0 phase. Its Hermite coefficient equals the normalized Gaussian integral and, by absolute Fubini, the finite auxiliary integral of that kernel. Both approximations

are disk-valued. The small-ball estimate on every fixed fibre and layer cake give softening L^2 error at most $(10/3)\eta$; truncation adds at most $e^{-\eta^2 T}$. Bessel transfers this error to any finite coefficient vector without a mode-count loss. The fixed parameters give radius at most $5 \cdot 10^{-15}$. The regular source budget still requires a coupled finite-integral-to-saved error ε_a at every address, weighted by the original Gaussian weights as $(5 \cdot 10^{-15} + \varepsilon_a)^2$. The exceptional squared term is added separately. A common coupled row radius $19/10^{11}$ suffices: the proved source mass is below $107/100$, so the existing regular and full source-error budgets follow by rational inequalities. No row certificate is inferred from this sufficient tolerance. Numerical certificates and full ϕ_T/P_3 response binding remain open.

The concrete regular-row certificate `Q0RegularSourceRow` fixes original ordinal 3818, panel 135 and the third node of its order-six rule. Its three saved binary64 values fit the sparse linear model with exactly zero residual. A centered check on the actual panel bounds the physical profile, giving whole-selected-vector Gaussian-to-saved radius below $2 \cdot 10^{-15}$, with omitted coefficients retained. Adding the finite auxiliary approximation gives radius below $7 \cdot 10^{-15}$ and meets the regular-row tolerance at this address. The byte slice, source metadata and row hash are pinned.

The near-zero source extension selects original ordinals 3469–4456, 10710–11238 and 17223–17542, containing 1,837 rows and 1,420 regular rows. Its compact payload preserves the original bytes and ledger hashes. Batched rational checks bound the fitted parameter by $2 \cdot 10^{-11}$ and the coupled model residual by $2 \cdot 10^{-22}$. Three actual panel blocks have scaled profile at most 10^{-10}. The shared analytic argument yields actual-to-saved and finite-to-saved radii below $6 \cdot 10^{-11}$ over the entire selected coefficient vector, inside the existing row allowance. The source-error, exceptional-profile, finite-Laplace and concrete row modules passed focused Lean CI at commit `a00103c2`, run `41725462291`. Repository-wide gates and independent review are separate delivery checks. The remaining rows and final upper bound remain work in progress.

The complete stored third Taylor center also has exact total and comparison-band energies checked by the Lean kernel. The proof assembles every grade's integer sum and divides by 2^{210}, then identifies the rational targets. The authoritative witness consists of ordered scaled-integer grade arrays regenerated from the same pinned binary64 input, preserving every coefficient. Its public energy and source-size theorems add no native axiom. The actual endpoint receives a fresh transitive audit after this replacement; the finite energies do not establish its Gaussian error hypotheses.

The endpoint proof is *work in progress*. One completed supporting calculation, `ProfileFirstGaussianCell`, proves $\int_0^1 (\frac{1}{2} - s)^2 e^{-s}\, ds = \frac{5}{4} - \frac{13}{4} e^{-1}$ and encloses it in $[0.054391816192, 0.054391816194]$ using proved exponential bounds. The complete profile and response estimates remain to be proved.

The pinned LeanCert dependency supplies a further completed supporting calculation in `LeanCertPilot`:

$$1 \le \int_0^1 \frac{2\exp(-u^2/(1+u^2))}{(1+u^2)^2}\, du \le \frac{5}{4}.$$

An explicit 16-cell dyadic certificate checks the enclosure and reciprocal domains with `decide +kernel`. A separate identity binds the integrand to the actual affine Gaussian phase integral for $u > 0$. Both theorems are audited without new native facts; the production Gaussian error estimates for the target endpoint remain work in progress.

The literal polynomial used in the q0 Wick calculation is linked to its actual parsed source by `Q0BinaryLiteralAssemblyBridge.sourceSum_eq_literal`:

$$\mathrm{Ssum}(\mathrm{q0RawWickEntries}; s, b, v) = \frac{\mathrm{he3}(\mathrm{sourceCoefficients}; s, b, v)}{60400508739434583230754315648473497600 0}.$$

This identity holds for all real s, b, v, ordered as squared radius, inverse damping rate and shifted argument. The source partition, chunk denominators, literal tables and recombination are checked in Lean; no source-table equality is assumed. A strict transitive audit allows only the standard

logical axioms. The inverse-rate factor already included in the literal is retained. Numerical cell certificates and the full upper-bound inputs remain separate obligations.

For the retained head through degree 501, `ProductionSmoothingApprox.production_smoothing_error` proves a relative L^2 error at most $4 \cdot 10^{-15}$ for the concrete rational smoothing matrices. The exact action in grade q is $S_{q+1} C_q S_q^{\mathsf{T}}$; orthogonality aggregates all 251 grades without multiplying the error allowance by their number. The entry rounding uses an integer square-root floor and only the standard logical axioms. The actual retained Gaussian coefficient enclosures and their complete energy-table comparison remain work in progress.

The actual Taylor sources also admit a proved finite auxiliary reduction in `ProductionFiniteLaplace`. With softening $\eta = 10^{-15}$ and cutoff $T = (2^{53})^2$, softening plus truncation costs less than $5 \cdot 10^{-15}$ in the full Gaussian L^2 norm. Smoothing and every Hermite-band projection preserve this bound. The regular auxiliary interval is $[0, 2^{53}]$, with $\eta^2 T \geq 40$. A LeanCert certificate proves $e^{-40} \leq 10^{-17}$ by kernel reduction. The finite Fubini identity binds the truncated response to the coefficient integrals; their numerical quadrature remains work in progress. `ProductionCompactLaplace` also restricts the first squared Gaussian radius to $[0, 80]$. Its radial error is at most 10^{-17}, and the combined softening and both truncation errors still fit the same $5 \cdot 10^{-15}$ reserve. `ProfileAdaptiveSource` permits the shorter regular auxiliary endpoint 100 on any certified radial region with $s \geq 1/16$ and $|\kappa(s)| \leq 1/100$. It retains 2^{53} elsewhere and proves that the combined source error for all selected regions still fits the original $5 \cdot 10^{-15}$ reserve, including after smoothing and band projection. The retained-cell certificates remain work in progress. The finite-union Gaussian pairing is identified with its rectangle integrals in `ProfileIntervalQuadrature`. Charge-one Parseval in `ProfileIntervalBands` then transfers retained coefficient-energy certificates to the original bands, without repeating the source reserve for each interval or coefficient. `ProductionRectangleAssembly` identifies those coefficients with a sum of retained nonoverlapping rectangles of the continuous unit-mode evaluator. An exact substitution $x = 1/u$ maps each auxiliary tail to a bounded interval with positive rational lower endpoint, including the Jacobian x^{-2}. No extra truncation or endpoint-zero premise is used. Finite partitions of these positive reciprocal intervals follow by mapping their boundaries back to the original coordinate. `ProductionMeshEnclosure` sums the certified regular and reciprocal cell radii, with the Gaussian normalization, into an enclosure of the actual retained coefficient. The common source reserve is applied afterward at the band level. The numerical cell and energy certificates remain open. The concrete partition in `SmallProfileMesh` uses seven radial strips. Kernel-reduced centered Horner checks establish the actual small-profile condition on three selected rational intervals. Thus its retained-source bound has no remaining profile-region premise; the numerical retained-energy certificates remain to be completed. Eight complementary centered-Horner rows in `SmallProfileSeparation` combine with the existing profile cover to give $|\kappa(s)| > 1/1000$ on the four other strips. Hence $\alpha x^2 + \kappa(s)^2 > 10^{-6}$ for $\alpha \geq 0$, uniformly even when the reciprocal coordinate is small.

The local affine-cell enclosure only requires continuity on the closed positive-coordinate rectangle. Coordinate clamping gives a continuous extension with the same integral. The complete normalized retained coefficients are real. Rational coefficient centers and radii give a finite sum of squared residual radii, which bounds the retained band by Parseval. The original source reserve is added once after this sum. The actual numerical enclosures and the endpoint remain work in progress.

The retained transfer is specialized to the fourteen actual captured bands. Their finite degree cutoffs and reserve inequalities are checked by kernel reduction. The source powers, stored centers and original caps are fixed; the remaining inputs are the numerical coefficient enclosures and their rational residual-energy checks.

Adaptive rectangular quadrature uses a binary tree of exact rational midpoint splits. Each leaf lies in the original closed rectangle, and the original iterated integral equals the sum of the ordered leaf integrals. The auxiliary coordinate remains outside the radial coordinate; continuity is required

only on the closed root. Actual numerical certificates for the complete leaf lists still have to be supplied.

Certified intervals for those leaves add to a root-integral interval by exact rational arithmetic. Pointwise analytic envelopes also give integral bounds by multiplying by the exact finite cell area. These operations preserve the separate source-error reserve and the original band budgets.

The fourteen regular and reciprocal roots are identified with the actual retained Gaussian coefficient, including the reciprocal Jacobian and local continuity. Rational bounds for $\sqrt{\pi}$ enclose its reciprocal, so interval multiplication accounts for normalization. Complete leaf certificates and rational endpoint checks then imply the coefficient bounds consumed by the original bands. The numerical leaf and energy certificates remain work in progress.

Separate rational intervals enclose π, $\sqrt{\pi}$ and $\sqrt{\sqrt{\pi}}$ with width at most 2^{-512} for profile normalization. The certificates use proved nested-square-root bounds for π and exact rational inequalities. Complete transformed node-shift intervals remain open.

Exact midpoint identities also combine already certified subtree intervals. The normalized coefficient can be assembled from the fourteen resulting root intervals, permitting separate proof batches without flattening the complete leaf list.

Both explicit base-pair integrals on the entire reciprocal root of radial strip 5 have rational enclosures with no numerical premises. A checked profile Horner trace and a LeanCert dyadic exponential certificate give pointwise envelopes; exact cell area and kernel-checked rational comparisons place the integrals in their original intervals. The first lower endpoint remains zero. The remaining roots and band-energy certificates are still work in progress.

The analytic-envelope leaf family reuses shared profile and scalar exponential certificates. Each leaf yields both actual iterated integral enclosures for arbitrary damping, preserving its original rational intervals. LeanCert dyadic checks and the finite rational comparisons use kernel reduction. On the reciprocal root of strip 3, three such leaves combine with an algebraic leaf. On the direct root of strip 5, twelve algebraic leaves combine. Exact midpoint identities prove both complete base-pair integrals from these checked cells without numerical premises. The direct roots of strips 1 and 3 similarly combine all twenty-four and fifteen algebraic leaves, respectively, with the same exact midpoint assembly. The explicit integrands are identified with the two actual production kernels at Taylor index zero, including the reciprocal Jacobian. These identities transfer the five complete-root enclosures at the fixed damping without numerical hypotheses. The remaining leaves, other roots and final band energies remain separate obligations. Conjugation invariance of the full Gaussian measure and actual base response also proves that both original Gaussian coefficients, including their rational frozen corrections, are real. Real interval errors thus give the same complex norm errors without a truncation premise.

The original Legendre brackets for orders 1 through 64 have exact integer sign certificates and Lean proofs of root existence and uniqueness. The actual conditional Gaussian coefficients satisfy a joint Bessel bound; their selected contraction indices and profile sign exponent $c+d$ are identified in Lean. Reflection is proved for the actual phase integrals and their normalized production coefficients, with the exact factor $\sqrt{(d+1)/c}$ for the lookup $(c,d) \mapsto (d+1, c-1)$ when $d+1 < c$. This factor is at most one; numerical reconstruction errors remain to be checked. A shared quadrature theorem derives the original weight formula and exactness below twice the rule order from finite polynomial identities, using integrated root quotients, Lagrange interpolation and polynomial division. Instantiating those identities for the original rules, the numerical weight bindings and analytic remainder remain work in progress. The semantic bridge now uses LeanCert's exact polynomial operations. Generated certificates for all 64 orders include the recurrence, orthogonality, derivative and Bezout identities, squared norms, and leading coefficients. Their full Lean replay remains work in progress. The family and its strict axiom audit are included in the main root so that the combined delivery checks all these obligations before acceptance.

A second route uses a finite-head estimate for the actual kernel directly; its coefficient enclosures and numerical head estimate remain work in progress.

Selected-spectrum comparison. The finite head only sees $S = \{(n, m) : n + 3m \le 501\}$. Write its approximation error as $G_T - X = \sum_r D_r$, where D_r is supported in degree band B_r with $\|D_r\|_2 \le e_r$. Let $w_p \in [0, 1]$ bound the retained absolute coefficient sum of the preprocessed block p. For nonnegative radii $A_r^2 \ge \sum_{p \in S \cap B_r} w_p^2 E_X(p)$, Lean proves

$$\sum_{p \in S} w_p |E_{G_T}(p) - E_X(p)| \le 2 \sum_r A_r e_r + \left(\sum_{g=0}^{3} \sqrt{\sum_{r \in I_g} e_r^2} \right)^2 .$$

Here I_0 contains the rounding remainder and I_1, I_2, I_3 are the three Taylor orders. Their bands are disjoint within each group, so their errors combine in quadrature before the groups are added. All pieces may share coordinates. Lean proves each coefficient cost by outward interval polynomial arithmetic; native arithmetic checks eighteen weighted radii from the stored table and its error allowance. The Gaussian component-error enclosures and their identification with the pieces remain explicit inputs.

Phase-source denominator. The parsed production phase has common denominator 3020025436971729161537 proved by bounded source-chunk computations and their assembly in Lean. This is source arithmetic; the phase cover and upper-bound Gaussian estimates remain separate obligations.

The conditional coefficients of the actual phase are defined in `F0ConditionalNode` and satisfy a coupled Bessel bound. For nonnegative radius and nonzero profile, `F0ConditionalEvaluator` identifies them with the signed historical shifted integral. `ConditionalNodeIndices` proves the exact selection of 31,877 distinct pairs from the larger 35,777-pair storage. The actual Gaussian integral reflection, proved by finite Stein/Ward recurrences and absolute Fubini, yields the production multiplier $\sqrt{(d+1)/c} \le 1$ when $d + 1 < c$. All these identities use only the standard logical axioms. Numerical node enclosures and floating arithmetic errors remain separate obligations.

The actual shifted coefficients $C_{cd}(A)$ also satisfy the charge relation

$$A(\sqrt{c+1}\, C_{c+1,d}(A) - \sqrt{d+1}\, C_{c,d+1}(A)) = (d + 1 - c) C_{cd}(A).$$

Finite shifted-moment identities and absolute Fubini establish this relation; Hermite normalization supplies the square-root factors. For $A > 0$ it reduces each row to the preceding row and the existing zero-row moment seeds, without a hypergeometric diagonal recurrence. A finite rational checkpoint checker encloses the actual positive moment ratios over parameter intervals, starting from the proved terminal bounds $[0, 1]$.

For the actual seed integral and every $x \ge 0$, LeanCert's exponential Taylor remainder on $[-x, 0]$ with derivative bound one gives

$$\left| J_0(x) - \sum_{n=0}^{N} \frac{(-x)^n J_n(0)}{n!} \right| \le \frac{x^{N+1} J_{N+1}(0)}{(N+1)!}.$$

Rounded integer checkpoints instantiate this estimate at the pinned row-227 rational midpoint squared. All 513 updates, the retained terms and the terminal remainder are checked, giving an actual integral enclosure of width $73351/2^{1024}$ without an assumed numerical accuracy. The negative term multiplier reverses interval endpoints; positive-scale soundness connects the integer guards to the integral. Concrete arithmetic uses `decide +kernel`, with reproducible hash-pinned inputs.

Common scalar and row-227 source binding. `ConditionalSeedNormalization` uses the positive actual seed $S(A) = (\sqrt{\pi}/2)AJ_0(A^2)$ for $A > 0$ to write $C_i(A) = S(A)U_i(A)$. The normalized coefficients have $U_{00} = 1$, retain the zero-row and charge recurrences, and satisfy $\sum_{i\in E} |U_i(A)|^2 \le S(A)^{-2}$ for every finite set of modes E. `ConditionalShiftEnclosure` and `ConditionalSourceRow227Shift` bind the original zero-based panel 4 / rule 31 address to source ordinal 227. Its profile is positive, its phase sign is one, and its actual coefficient vector equals the canonical vector at the physical shift A_{227}. With m the pinned rational midpoint,

$$|A_{227} - m| \le 2 \cdot 10^{-18}.$$

`ConditionalSourceRow227Scalar` combines the actual midpoint moment enclosure and square-root-of-π bounds to give

$$S(m) > \frac{49}{50}, \qquad |S(m) - \widehat{S}| \le 784 \cdot 10^{-24},$$

where $\widehat{S}$ is the recorded exact dyadic center. No numerical-accuracy premise is used, and m is not identified with the physical shift. Midpoint-to-physical-shift transport, certified normalized-triangle arithmetic, saved-row coefficient errors and final energy-table agreement remain separate obligations; these results do not establish the final upper endpoint.

`ConditionalReconstruction` proves that nonconstant coefficients vanish at a zero profile, retaining the constant phase $\mathrm{ph}(\sqrt{s})$. Reflection therefore extends to all nonnegative radii. The exact completion of the selected 31,877 mathematical modes has weighted error at most 10/7 times the coupled canonical error at the same nodes and weights. This consumes explicit canonical error and nonnegative-weight premises; physical quadrature, primary contraction, floating operations and table agreement remain separate obligations.

Exact phase moments. For $\varphi = De^{-\varepsilon(U+V)}S(U,V,W)$, clear the rational coefficients as $A = LS$. Each even moment is a finite rational sum over the coefficients of A^{2j}, divided by L^{2j}, at Gaussian rate $1 + 2j\varepsilon$. The integral identity applies directly to the stored phase and does not use the numerical phase-cover certificate. Separate native checks give exact rational values for $\mathbb{E}\varphi^2$, $\mathbb{E}\varphi^4$, and $\mathbb{E}\varphi^6$, approximately $2.56852947444587 \cdot 10^{-6}$, $5.71439509703932 \cdot 10^{-10}$, and $1.68238736430120 \cdot 10^{-11}$. These equal the squared norms of the three positive Taylor sources. The same finite evaluator with a restricted Gaussian measure also proves $\int_{|z_1|^2>50} \varphi^2 \, d\gamma_2 < (5 \cdot 10^{-15})^2$. A native check bounds the exact rational factor after extracting e^{-155}; the Taylor polynomial through degree twelve at 5 gives $e^5 > 148$, hence $e^{-155} < 148^{-31}$. This certifies the radial tail of the production phase. Since $|\varphi| < 1$, every positive Taylor source restricted to this event also has $L^2(\gamma_2)$ norm below $5 \cdot 10^{-15}$. Sharp comparison-band errors still need the captured Hermite energy and approximation bounds; the total moments alone do not change the upper bound. The orthogonal tail interface subtracts the separate captured energies $\sum_i \max(0, \ell_i - e_i)^2$ from the total moment, retaining each band's error cap. An alternative interface uses one real inner product and the approximant norm to bound the whole approximation error. The remaining numerical coefficient and integral bounds are explicit hypotheses. Finite normalized Hermite expansions have exact rational squared norms. The preserved coefficient inputs supply a complete finite check for the third approximant: all 6,552 coefficients through degree 51, its total norm, and its four projected band norms. Injective graded indexing also proves that its higher-degree tail vanishes. Its Gaussian approximation error remains a separate obligation; the stored error arrays are not treated as proofs. A continuous finite-polynomial representative satisfies the charge-one function interface. The third-order tail cap follows from four explicit finite-band Gaussian error caps by subtracting captured energy from the exact sixth moment; the Gaussian caps themselves remain hypotheses. The finite rounding difference $c_1 + c_2/2 + c_3/6 - v$ has norm at most 10^{-18}, proved by exact integer squares at a common dyadic denominator. The three other preserved center sources and retained sizes are checked separately from Gaussian approximation errors. The frozen witness

uses typed integer grades at scale 2^{206}. Standard-axiom proofs establish all grade sizes and the complete production coefficient-square sums for grades 0 and 90; these component energies do not discharge the mixed Gaussian energy-table comparison. The first and second centers also have exact certificates for their six and four captured bands. The concrete conditional theorem fixes every center and its polynomial representative and derives the rounding and all three omitted tails internally. Its remaining inputs are fourteen explicit finite-band Gaussian error inequalities and the concrete energy-table accuracy. The slightly tighter finite error caps are generated from the original allowances; their Gaussian validity remains open. The weighted Gaussian kernel has a common checked parameter predicate. It gives real covariance positivity, genuine integrability and the uninserted physical-pair mass ν^2/δ. A product-measure argument combines two pairs into $\nu^4/(\delta_0\delta_1)$. One rational parameter table covers the fourteen selected finite bands and four coordinate exclusions. These parameter checks do not assume the numerical tail bounds recorded in the input files. The actual Hermite coefficient identity and the weighted matrix-energy reduction are supplied by the coherent-transform and rectangle-operator constructions below. Numerical tail estimates remain separate obligations. A Gaussian Laplace representation of complex phase is proved exactly, including the zero argument, integrability and a uniform absolute-integral bound. Absolute Fubini exchange applies this representation to the actual Taylor inner products. Affine completion of squares and translated Gaussian Wick sums give exact conditional monomial integrals with the inverse-rate mass. Finite support extends this evaluation to the same complex polynomials used by the Hermite basis, with integrability proved term by term. For any positive radial rate r and fixed complex source a, translation by $\overline{a}/r$ evaluates finite polynomials in z and $\overline{z}$ under the tilt $\exp(-(r-1)|z|^2+2\,\mathrm{Re}(az))$ through shifted Wick sums. The factor $\exp(|a|^2/r)$ and the Gaussian mass $1/r$ are retained. A half-rate majorant proves integrability of these polynomial insertions. For arbitrary independent complex sources a, b, a half-rate majorant is uniform on bounded source sets for $\exp(-(r-1)|z|^2+az+b\overline{z})$. It proves absolute integrability of every mixed monomial and justifies complex differentiation under the integral. Differentiation in a inserts z, and differentiation in b inserts $\overline{z}$. For fixed r, b, define the finite polynomial recurrence

$$P_{0,q}(X) = r^{-q}X^q, \qquad P_{p+1,q}(X) = P'_{p,q}(X) + (b/r)P_{p,q}(X).$$

The exact (p, q) moment is $r^{-1}e^{ab/r}P_{p,q}(a)$, and finite linearity evaluates every finite polynomial insertion. The mass and means are respectively $r^{-1}e^{ab/r}$, b/r, and a/r. Applying this recurrence to the actual multiplier $(i\varphi)^j$, excluding F_0, gives its coherent transform with no source-identity hypothesis. The rational phase polynomial uses the production term list in physical order $(z, \overline{z}, w, \overline{w})$, and the radial rate is $r = 1 + j\varepsilon$. For each physical coordinate, outgoing sources (x, y) and incoming sources (u, v) have factor $\exp(x\overline{z}+yz-xy)\exp(uz+v\overline{z}-uv)$. All eight independent sources admit absolute integrability, and the transform is a finite polynomial times the product of two factors

$$r^{-1}\exp\bigl(-(1-r^{-1})(xy+uv)+r^{-1}(xu+yv)\bigr).$$

The total normalization is r^{-2} under the normalized Gaussian product measure. For the real bipartite coupling matrix M, the existing gate $I-M^TM>0$ gives a positive rate $\epsilon = \min(1/2, \kappa/3)$ by completion of squares, where $\kappa = \det(I-M^TM)/\operatorname{tr}(I-M^TM)$. Every finite polynomial times the coupled Gaussian is consequently in L^2. A polynomial-coefficient lift of the moment recurrence identifies the actual scaled transform with this finite-polynomial kernel, retaining both mass factors. The normalized holomorphic monomials in eight Gaussian sources are orthonormal, so their coefficient energies are summable and bounded by the squared source norm. Gate coordinates are $((y, u), (x, v))$ per physical coordinate; matrix indices (m, n) therefore map to $((m_q, n_p), (m_p, n_q))$ explicitly. The actual normalized weighted Hermite matrix identity has complete source proofs in the gate-jet extension below. Exact binding to the historical Laguerre polynomial and numerical Hilbert–Schmidt caps remain open. The finite coefficient construction evaluates finite physical Hermite generating blocks and recovers exactly their retained normalized matrix entries. Separately, each coefficient of the full scaled coherent transform has a finite factorial

formula: source-space L^2 justifies Fubini, and the unit-rate moment $q(q-1)\cdots(q-p+1)a^{q-p}$ for $p \le q$, with value zero for $p > q$, extracts the first source pair, leaving a polynomial Gaussian moment in the second. These are total finite formulas. Every actual physical matrix entry also has a finite damped Wick formula, with diagonal moment $p!/r^{p+1}$ and zero off the diagonal. The ground coefficient is identified with the actual ground matrix entry. The gate-jet extension below supplies complete source proofs of the weighted Hermite identity for all indices; numerical caps remain open. Finite Gaussian source extraction equals a mixed derivative at zero. In the fixed gate order $((y,u),(x,v))$, the same finite derivatives are evaluated by the physical Hermite contractions, without commuting derivatives or exchanging physical and source integration. Independent diagonal scaling and tensoring then identify every coefficient of the actual scaled transform with a finite sum of physical Hermite moments from the production polynomial, retaining both mass factors and the explicit index permutation. Separate finite identities identify the Hermite normalization and show that the squared source scales equal the original energy weights. Exact factorial and scale identities, together with finite physical polynomial integration, complete the source coefficient equality with the normalized weighted matrix entry. The full weighted energy is summable, and its full and finite sums are bounded by the Gaussian source integral of the squared norm of the actual scaled transform. Numerical bounds on that source integral, Hilbert–Schmidt caps, the historical Laguerre binding, and the final upper endpoint remain open. The radial normalization is identified directly with the production Hermite polynomial: for $s \ge 0$,

$$h_{m+k,m}(\sqrt{s}) = (-1)^m(\sqrt{s})^k\sqrt{\frac{m!}{(m+k)!}}\,L_m^k(s).$$

The value is real, so conjugation introduces no additional sign. The natural power of the radius includes zero radius and zero charge. The conditional block $d \le q < 251$ is reindexed by $m = q-d$; its charge $k = d+1-c$ lies in $\{0,\dots,251\}$ and its length is at most $251-(k-1)$, where subtraction is truncated at zero. Both charges zero and one have maximal length 251. These finite identities and index maps do not discharge a numerical radial Gram bound, quadrature error or the final endpoint.

The Gram construction expands the actual source norm into finite Gram sums using the production polynomial's support and coefficients. It preserves all cross terms. Each pair Gram entry is evaluated by shifted Gaussian moments followed by a finite correlated polynomial moment; the original Schur conditions and both kernel masses remain explicit. Substitution of the actual gate parameters gives a finite scalar expression for the squared source norm, which bounds the full actual weighted matrix energy. Its numerical enclosure and the final endpoint remain open. The rectangle-tail construction retains the exact coordinate cutoff factor $R^D T^{L+1}$ and covers the complement of the inclusive input rectangle $(500,500,251,250)$ by four coordinate exclusions. The existing comparison-band definition gives the outgoing degree bound for each finite band. All fourteen finite, nonzero-power bands are bound to consecutive quartets of the original 56 checked gates. Each actual rectangle-tail energy is summable and bounded by the sum of the four finite source-norm expressions divided by their cutoff weights. A finite complex Cauchy–Schwarz argument bounds each selected-band matrix action by this same sum times the input coefficient energy, without an extra mode count. Finite Gaussian integral exchange identifies the action with output coefficients of the actual Taylor multiplier applied to a finite normalized Hermite input. Under explicit `PhiBound`, the production multiplier defines a continuous operator on physical L^2. Its matrix is identified with the actual Gaussian entries, so completeness of the Hermite basis gives convergence of the finite rectangle-tail actions. Finite Bessel and output Parseval then extend the same four-source bound to every physical L^2 input and the full selected output band. Numerical source bounds and the final closed endpoint remain open. For the correlated form $Q = A|z|^2 + 2B\,\mathrm{Re}(z\overline{w}) + C|w|^2$, the exact conditions $A > 0$ and $AC - B^2 > 0$ give the positive coercive rate $(AC-B^2)/(A+C)$. A separable Gaussian majorant then proves joint integrability of every finite mixed polynomial. Fubini and the first-coordinate mean $-(B/A)w$ reduce its integral

to a finite Wick sum at outer rate $(AC - B^2)/A$, with zeroth mass $1/(AC - B^2)$. Full bipartite polynomial insertion belongs to the source-space batch above; the historical operator identity and numerical tail estimates remain open. For each finite center, fixing the first coordinate gives its exact polynomial slice in the second coordinate. The actual source $\overline{c}(i\varphi)^j$ factors into this polynomial and Gaussian damping with conditional rate $\alpha = 1 + j\varepsilon$, giving an exact conditional finite Wick sum. Fixed-parameter boundedness justifies Gaussian product Fubini, and the actual inner product becomes the auxiliary and first-coordinate Gaussian integral of this finite evaluator, with no numerical integration premise. For each unit Hermite mode, a rational Wick polynomial gives the exact regular quadrature kernel in squared-radius form. A pointwise polynomial enclosure on a normalized square implies a bound for the real production cell integral and shows its imaginary part vanishes. Shared moment recurrences agree with the rational Wick compiler for positive inverse rate; finite table reuse requires proved base and successor equations. The real moments and their affine combination admit inverse-rate factorizations, and reciprocal-variable rate identities remain valid when the profile vanishes. These identities provide no numerical enclosure. Smooth angular moments

$$M_h(x) = \int_0^{\pi/2} \sin^{2h}\theta \, \cos^2\theta \, e^{-x\sin^2\theta} \, d\theta$$

are positive, have consecutive ratios in $(0, 1]$, and satisfy the exact three-term recurrence. For $x \geq 0$ their ratio update equals a monotone abstract update, for which finite backward interval propagation is proved from a terminal interval in $[0, 1]$. The angular ratio is bounded by the iterations with terminal values zero and one. The original moment integrand is interval integrable even at order zero, and the continuity, derivative, range and open-interval pullback identity for the angular substitution are proved. The integral change of variables gives $J_h(x) = (4/\pi)M_h(x)$, with endpoint equality handled almost everywhere. The original moments inherit positivity, the recurrence, and finite backward ratio enclosures for $x \geq 0$. The substitution $v = t/(1 + t)$ gives the integrable real half-line kernel identity

$$\int_0^\infty \left(\frac{t}{1+t}\right)^h t^{-1/2} \frac{e^{-xt/(1+t)}}{(1+t)^2} \, dt = \frac{\pi}{2} J_h(x).$$

Absolute Fubini and the exact shifted Gaussian moment identify the normalized zero-row coefficient with

$$I_{0,d}(A) = \frac{\sqrt{\pi}}{2} \, \frac{A(-A)^d J_d(A^2)}{\sqrt{d!}}.$$

This is connected to the actual canonical evaluator and production coefficient, including the profile sign to power d. At $d = 0$ that factor is one. These identities use only Lean's standard axioms. The exact zero-row recurrence $I_{0,d+1}(A) = -Ar_d(A^2)I_{0,d}(A)/\sqrt{d+1}$, which is proved for the actual Gaussian coefficient and transferred to the signed production value. The canonical coefficient has norm at most one by Cauchy–Schwarz in Gaussian L^2. Its two-term raising sum has norm at most $\sqrt{c+1} + \sqrt{d+1} < 35$ on the stored index rectangle. The parameter derivative identity and midpoint transfer still require proof. Conjugation invariance of the Gaussian and rational Hermite coefficients prove reality of every canonical coefficient and the corresponding signed production coefficient. Their complex distance to a real stored center equals the absolute real-part error. The saved-center interface adds half the physical-parameter interval error and the explicit model-to-saved error to $5\rho^2$; the displayed rational budget fits below 10^{-12}. Actual node and byte certificates are still required. The corresponding actual Gaussian approximation has L^2 and all-mode coefficient error at most $5\rho^2$, where $\rho = \sqrt{1 + a^2}v(s)/\sqrt{s}$ for $s > 0$. Saved-entry enclosures remain separate. The constant-plus-linear phase approximation has pointwise error at most $\frac{5}{2}\|h\|^2/r^2$ at reference radius $r > 0$, also when $r + h = 0$. Its actual Gaussian coefficient and node-data instantiation remains separate. The zero-safe inequality $\|\operatorname{phase}(r+h) - 1\| \leq 2\|h\|/r$ for $r > 0$ bounds the actual conditional phase and coefficient error by $2\sqrt{1 + a^2}|v(s)|/\sqrt{s}$ in Gaussian L^2, including at profile zeros. Its use in a final node and energy ledger remains to be proved.

Numerical node certificates and the complete upper bound remain open. The cell enclosure is an explicit hypothesis; no numerical cell certificate or numerical improper-integral bound is claimed. The recurrence $M_{p+1,q} = \mu M_{p,q} + (q/\beta)M_{p,q-1}$ preserves the affine cancellation for every mixed degree, including zero mean and $q = 0$. Real radial powers with exponent greater than -2 also have their exact damped Gamma integrals and integrability proved. The zeroth conditional bracket simplifies exactly to $\alpha z/\beta^2$, where $\beta = \alpha + a^2 t$. For $|z| \geq r_0 > 0$, its conditional-kernel norm has the profile-uniform majorant $2t^{-5/2}/(\alpha r_0^3)$ and an integrated tail above $T > 0$ bounded by $4T^{-3/2}/(3\alpha r_0^3)$, including when $a = 0$. The full polynomial-weighted source tails, the radial endpoint and numerical integration enclosures remain open.

Radial repair has a quantitative ceiling. Lean certifies the polynomial disk map $z \mapsto zS_6(1 - |z|^2)$, where $S_6(q) = \sum_{j=0}^{6} \binom{2j}{j} q^j/4^j$: the exact defect $1 - (1 - q)S_6(q)^2$ is q^7 times a polynomial with positive coefficients. Applied to $P_3(\varphi)$ this is a degree-39 factor. But Lean also proves $(1 - |P_3(x)|)^2 \leq x^6/1225$ for $x^2 \leq 1$. For charge-one responses and any same feasible preprocessors, the full coefficient change obeys

$$\|K_H - K_G\|_A \leq (\|H\|_2 + \|G\|_2)\|H - G\|_2.$$

Thus outward radial changes of disk responses cost at most $2\sqrt{\mathbb{E}\varphi^6/1225}$ in the full kernel. The integrated Lean theorem retains the sixth-moment and radial-gap premises explicitly; it does not identify the concrete Gaussian moment or certify a new score. These results leave the endpoints unchanged. Preprocessor changes must be compared at the same P, Q; this ceiling does not bound a separate gain from reoptimising them.

2.2 Lower bounds: Hermite projection games

A lower bound needs an operator T whose *vector* value is close to 1 while its *scalar* value is small. We use

$$T = P_{1,0} - \epsilon\, P_{2,1}, \qquad \epsilon = \tfrac{149}{200},$$

the projection onto linear functions minus a multiple of the projection onto the cubic chaos.

- *Vector value.* With $u_\ell = z_\ell/\|z\|$, $\sum_\ell \langle u_\ell, Tu_\ell\rangle \geq 1 - (1 + \epsilon)/(n + 1) \to 1$.
- *Scalar value.* For all $|f|, |g| \leq 1$, $|\langle f, Tg\rangle| \leq B = 0.713420886479$.

A finite-partition transfer turns this into matrices, so $K_G^{\mathbb{C}} \geq 1/B > 1.401697117301$. The complete cover has 10,909 boxes. Two added pure-row law nodes and three mixed-source Young nodes use the original source caps and Stein cuts, so the Gaussian assumptions are unchanged. Their batches, every box and the full partition are checked separately. Adaptive binary splits reprice the existing components while preserving their scale interpolation domains. All generation inputs are retained in `data/lower-unshifted/`.

The scalar cap is the hard part. Writing $h = \frac{1}{2}(f + g)$, $k = \frac{1}{2}(f - g)$,

$$\Re\langle f, Tg\rangle = t^2 - \epsilon\|P_{2,1}h\|^2 - \|P_{1,0}k\|^2 + \epsilon\|P_{2,1}k\|^2, \qquad t = \|P_{1,0}h\|.$$

Only the last term is positive: we must show that a *bounded* function cannot have too large a cubic part without paying elsewhere. After a rotation $P_{1,0}h = t\, z_1$, and the cubic part of k splits into pieces that involve z_1 (*visible* sectors) and a *remote* residual $q \in H_{2,1}$ in the other coordinates. The proof then has three layers.

1. **Row reduction.** For each of 19 profiles r, the contribution of the visible sectors is the supremum of an explicit finite-dimensional Gaussian integral, a convex one-variable *potential* Ψ_r, evaluated at the scale of the remote residual. One gate variable is shared across rows.

2. **The law of the remote cubic.** What remains is a statement about the distribution of $U = |q|^2/\mathbb{E}|q|^2$ for an arbitrary cubic Wick polynomial in arbitrarily many variables. We use exact identities and inequalities valid for *every* such cubic: the sharp fourth-moment cap $\mathbb{E}U^2 \le 44$; the sixth-moment bound $\mathbb{E}U^3 \le 9m - 12 + 40.45313\,(m-2)^{3/2}$, $m = \mathbb{E}U^2$, which rests on thirteen tensor-contraction facts including the Wick expansion of the sixth cumulant; a first-moment bound; and families of *Stein cuts* $\mathbb{E}\,G(U) \ge 0$, obtained by Gaussian integration by parts and completing squares against spectral budgets, a covariance order and *contraction-Gram budgets* (positive semidefinite Gram matrices of the coefficient tensor).

3. **A finite certificate.** The parameters are covered by 10,909 boxes in (t, τ, x); on each box a positive combination of rows and node certificates (351 existing mixed nodes, two added pure-row law nodes, three added mixed-source Young nodes and 9,237 row nodes, each a polynomial majorant checked in exact rational arithmetic) bounds the scalar value by B.

2.3 The primal/dual viewpoint

For the real constant, Heilman, Jones and Malavolta proved an exact minimax: over rounding schemes μ with kernel coefficients $c_m(\mu)$ and Hermite projection games $Q = \sum_m q_m \Pi_m$,

$$\frac{1}{K_G} = \sup_{\mu} \Big[c_1(\mu) - \sum_{m \ge 3} |c_m(\mu)| \Big] = \inf_{Q} \ \mathrm{OPT}(Q),$$

The complex programme uses charge-one bidegrees. A global complex minimax identity is not established here; the two endpoint proofs use separate one-sided transfer theorems. An upper bound comes from a realizable kernel, and a lower bound from a game with a uniform scalar norm cap. Finite-pool duality alone does not certify either global claim.

3 What exactly the computer checks

Proved in Lean	**Assumed (certified outside Lean)**
Lower bound. `lower_bound_of_refined_cover_data : NodeCaps → IsGrothendieckBound K → 1.401697117301 < K` the witness T on L^2 of the Gaussian measure and its vector value; the transfer to matrices; the polarisation identity; the per-row Gaussian reduction of $\langle f, Tg\rangle$ for every unimodular pair (frames, fibre decomposition, Parseval, gates); all moment, Stein, spectral and contraction-Gram inequalities for every circular cubic, including the sixth-cumulant Wick identity; the finite certificate (five row, three original-batch, one added law-batch and five refined-cover native checks)	`NodeCaps`: the 19 potentials Ψ_r (defined in Lean) lie below their stored rational caps at the tabulated nodes. `lower_bound_of_capsRest` proves the same endpoint from `CapsRest`: the vertex caps of rows 12–15 ($t = 0.84$–0.855) and any excluded-node caps, using the same refined cover. The 60 range checks for fifteen rows use a computation-only layer. CI replays the final assembly when the 53-file cache provenance matches; a mismatch triggers the runner build and the axiom audit. Four rows retain external vertex caps.
Upper bound. `UpperChaos.fresh_upper_rounder`, for an L^2 charge-one truncation X under component-error and table-accuracy inputs, proves $K_G^{\mathbb{C}} \le 50000000000000000/355973699960428191$. the concrete disk-valued rounding; tensor Hermite completeness by Fubini, Gaussian mixing by characteristic functions, Parseval and Mehler, charge selection, summability; the transfer inequality (tensor powers); kernel composition and locality; the head through degree 501 from 21,168 stored energies by a verified interval evaluator (one `native_decide`); the tail bound; the endpoint; the selected-block comparison and native band masses	`ComparisonInputs`: an exact error decomposition, spectral supports and inherited Gaussian component norm caps. The concrete Taylor specialization leaves fourteen finite Gaussian band-error estimates and the energy-table discrepancy as hypotheses. Lean derives the selected comparison $851922301/10^{15}$ (93.57% of the residual allowance). The decomposition itself is explicit in Lean (`difference_pieces`, `fresh_upper_pieces`): 18 band pieces of the smoothed Taylor errors and the rounding remainder, so only their 18 norm bounds remain external. Finite table accuracy $6847009/(5 \cdot 10^{17})$. The Hardy tail $303362486/10^{16}$ (3.33%, weighted square sum $\le 1737972/1000000$ at radius $4149/4000$) is proved (`HardyReal.hardy_GT_tight`). The native head is $28153/10^{12}$ (3.09%).

The degree-531 canonical pair table indexes conditional Hermite node values by pairs (c, d) satisfying $0 \le c \le 266$, $0 \le d \le 265$, $c \le d + 1$ and $c + d \le 531$, ordered by $(c + d, c)$. Lean verifies its finite rank, ordinal coverage and uniqueness certificate. No downstream Lean result yet consumes this certificate; the Gaussian node-value enclosures remain part of the upper-bound work in progress.

The repository-wide audit also covers the native facts pinned in `data/axioms/native.txt`, read by `scripts/check_axioms.py`. These check the upper head, selected-band masses, the six-part phase cover, the second, fourth and sixth phase moments and the squared radial phase tail, cap chords, the lower row/batch certificates, the additional pure-row and mixed-source law nodes and the complete finite lower cover. The native facts (cap-range checks, Hardy region and circle checks, and others) are listed in `data/axioms/native.txt`; the report is `scratch/Axioms.lean`, with its required declarations in `data/axioms/required.txt`. No `sorry`; normal `ready/*` delivery CI builds the single target `Grothe`, cap rows included, and checks its one axiom report before merging. The cap-range checks take about 21 CPU-hours cold, so CI replays them from a verified build cache. Its provenance hash covers 48 local computation modules and five directly hashed

extras (53 files total); the extras' imports are not traversed. The numerical checker has its own precompiled library, `GrotheCapsKernel`, whose constants retain the `GrotheKernel` namespace. The caps digest is part of the build-cache key: when no verified cache matches, CI runs the same check on the self-hosted runner, which builds, audits and marks `.lake/build/caps-closure.sha` before saving the cache. That persistent build supplies later CI replay; no separate manual droplet verification is needed. A digest alone is not verification. The localised-piece checkers likewise import only computable definitions and rational data; their campaign and final wiring remain separate. The generic response and quintic-transfer lemmas add no native facts or endpoint assumptions. The lower bound still needs numerical caps; the sharp upper bound retains external Gaussian component-error and table-accuracy inputs. The four existing arc checks each give at most 273/100; the proved bound on π turns their sum into the circle-mean budget 1737972/1000000. Before the tail Cauchy–Schwarz step, the known linear coefficient's weighted square is deducted from that budget. The abstract table interface uses this sharper Hardy input; the concrete rounding proves it without new Gaussian hypotheses. Holomorphy under positive kernel decay, odd-series continuation and circle Parseval are proved. The holomorphy and Hardy assembly also cover separate measurable disk-valued left and right responses (`HholPair`, `hardy_of_pair_inputs`); the incumbent results are specializations. The decay region and the circle mean are checked in Lean by five interval facts. With them, `hardy_of_real` reduces `hardy` to summability and the real Gaussian identification of the odd series, which `HardyReal.hreal_GT` proves for G_T: the smoothed Mehler pair has the law of $(U, RU + SW)$ (characteristic functions), whose density on $\mathbb{R}^8$ is the holomorphic kernel at the real cross covariance (Cholesky substitution). Hence `hardy_GT`: the Hardy input is proved. The research modules also prove 13 algebraic facts for coupled complex modulus majorizers, finite descent and semiconvex chord assembly, using only the three standard axioms. Their Gaussian integration, response extraction and global-cover obligations remain separate; these facts change neither endpoint nor its hypotheses.

4 Discoveries along the way

Status labels: [**Lean**] machine-checked in this repository; [**P**] written proof; [**P,C**] proof with an exact or interval-arithmetic certificate replayed independently; [**F**] floating-point evidence only.

Other upper-bound variants and their verification records are preserved outside the active build in `miscellaneous/other-bounds/`. Shared Gaussian phase identities and Laplace representations remain part of the sharp target.

4.1 Common responses, common tails, and an architecture separator

[**P; supporting algebra in Lean**] A symmetry-closed family of actual cubic responses has an exact five-mode first/cubic image and a fixed-size original-noise replica representation. It preserves the shared quadratic drivers: a scalar Gaussian used to describe a cubic's law must retain third-chaos noise correlation ρ^3, rather than acquire first-chaos correlation ρ.

[**Lean**] All correlation samples of a response pair use the same signed Hermite coefficients. Prices a_i with $|\sum_i a_i \rho_i^{2k+1}| \le 1$ for all $k \ge 2$ bound the complete omitted absolute tail. A finite front and one geometric remainder budget certify every index. The resulting separator is linear in response columns, so it bounds whole convex pools, with explicit integration-error costs. Numerical pool screens still require certified input integrals; a favorable relaxation value is not an achievable rounding score.

[**Lean**] An exchange can retain the incumbent's signed coefficients with one error allowance. If $\sum_k |c_k - \widehat{c}_k| \le \xi$, then every admissible spectral price has error at most ξ, and the full reciprocal score satisfies $D(c) \ge D(\widehat{c}) - \xi$. The same coefficient error must not be charged independently at every sampled correlation. The current bounded phase-perturbation searches use this interface

for discovery; their Gaussian integrals still require certification before they can improve either endpoint.

[P,C; polynomial certificate in Lean] The entire Gaussian-phase cubic-preprocessor family $H_{\mathrm{phys}}(A\rho + B\rho^3)$, $|A| + |B| \le 1$, including arbitrary complex phases and convex mixtures, has full reciprocal score at most 0.71186. This is below the incumbent $0.7119473997\ldots$. The separator is $\Re(c_0 - \frac{9}{25}c_1 - c_2 - c_3)$; positive polynomial certificates handle all complex phases. Thus optimizing this class further cannot improve the current upper bound. General non-Gaussian cubic responses and nonlinear shared image fields are not excluded. The Gaussian-family coefficient extraction is a written application of the Gaussian phase identity; it is not yet connected end to end in Lean.

Phase errors and exact symmetries [Lean]. Phase alignment has an exact modulus-weighted quadratic defect. An unweighted degree-moment budget is unavailable: the phase coefficients satisfy $r_k \ge 1/(4k(k+1))$ for $k \ge 1$, so $\sum_k (2k+1)c_k = \infty$. Certificate comparison can instead align an anchor by a common isometry: if the measured blocks are transported with it and the remaining error is at most e, the complete energy-spectrum discrepancy is at most $2\|x\|e + e^2$. A pure symmetry has zero price, including its entire tail. The kernel is preserved only when its correlation operator is transported too; rotating only the response can change the anisotropic kernel used here. An additional polarized certificate prices the whole omitted energy change by AB, where A, B bound the approximation errors of the response difference and sum by retained polynomials. These may be certified with positive squared-error integrals, avoiding subtraction of nearly equal moments in a residual Gram matrix. Two overlapping coordinate tests either keep the phase argument away from zero or resolve it by an affine polar chart. Lean checks the exact transformed Gaussian density and a Bernstein range enclosure. A symmetry-invariant cutoff preserves admissibility and costs at most its omitted mass in squared L^2 distance. This lets an upper construction use compact chart domains with an explicit full-spectrum allowance. An exact rational phase update has unit modulus for every real field, with a cancellation-free increment formula. A finite radius jet retains the radius and angle as independent polynomial variables until cancellations finish. Lean checks carry-free coefficient indexing, exact translation, and the finite remainder bound on complex patches. A norm-only homogeneous transport certificate also has a proved ceiling: its certified score cannot exceed b/N^2 when the anchor has norm $N > 0$ and score at least $0 \le b \le N^2$. This limits that certificate formula, not the corrected response's true score. The update assembly keeps the signed retained changes, passes the positive error price through contractive smoothing and the bivariate kernel, and connects the entire score to the transfer theorem. The reference error is the inherited anchor error plus finite weighted retained-data errors, charged once. These forty-seven Lean results change neither endpoint. The compact Gaussian integrals still require certification.

4.2 Inequalities for cubic Gaussian chaos

Throughout, q is a circular cubic Wick polynomial (bidegree $(2,1)$) in any number of complex Gaussian variables, normalised by $\mathbb{E}|q|^2 = 1$, and $U = |q|^2$.

- **Sharp fourth moment [Lean].** Every real third-chaos variable satisfies $\mathbb{E}F^4 \le 93\,(\mathbb{E}F^2)^2$, with equality for $x^3 - 3x$. For the cubic kernel, write $N = \sum f^2$, $s = \|f \star f\|^2$ and $h = \|\operatorname{Sym}(f \star f)\|^2$. The Wick identity $\mathbb{E}F^4 - 3(\mathbb{E}F^2)^2 = 1296(s + 3h/2)$, together with $\mathbb{E}F^2 = 6N$ and $h \le s \le N^2$, proves the bound directly, without a separate positivity premise on the Gaussian expectation. For a circular cubic, the complex contraction identities and Hilbert–Schmidt estimates instead give $\mathbb{E}U^2 \le 44$ at any dimension. Equality is attained by $Z(|Z|^2 - 2)/\sqrt{2}$ for a standard complex Gaussian Z; the equality calculation is analytic. The Lean theorem is `Circular.fourth_moment_le_variance`. It implies $x \le \sqrt{42}$, while the existing certificate still uses $x \le 31/4$.

- **Sixth moment [Lean].** $\mathbb{E}U^3 \le 9m - 12 + C_6(m-2)^{3/2}$ with $m = \mathbb{E}U^2$ and $C_6 = 40.45313$. The Wick monomial mixed with Gaussian noise attains 40.295, so the constant is sharp to 0.4%. The inequality is *false* without circularity: two real cubics $He_3(x_0)/\sqrt{12}$, $He_3(x_1)/\sqrt{12}$ give $m = 47$ and $\mathbb{E}U^3 = 13506 > 12622.55$.
- **The sixth cumulant as a sum over graphs [Lean].** The Wick identity for κ_6 of a cubic chaos is proved by expanding $\mathbb{E}F^2, \mathbb{E}F^4, \mathbb{E}F^6$ over closed cubic multigraphs, one Gaussian integration by parts per edge ($69 + 15 + 3$ certified steps). The disconnected graphs cancel, and the six connected cubic multigraphs on six vertices give the tensor contractions.
- **Contraction-Gram budgets [Lean].** A new family of universal inequalities between the chaos components of U and of q^2. With $X = \Gamma(q,\bar{q})/3 - 1$, $c = \Gamma(q,q)/3$, $r = \Re(c\bar{q}^2)/U$, $P = \mathbb{E}[XU]$, $Q = \mathbb{E}[rU]$, for example

  $$\mathbb{E}X^2 + \tfrac{1}{10}\mathbb{E}r^2 \le \tfrac{34}{45}P - \tfrac{1}{18}Q.$$

  Each budget is certified by two positive semidefinite 2×2 Gram matrices of contractions of the coefficient tensor.
- **First moment [P,C].** $\mathbb{E}|q| \le 0.97073$, from a two-plane small-ball estimate (a winding argument on S^3 and a horizontal Lipschitz bound). The sharp value $\sqrt{\pi}/2 = 0.8862$, which the Gaussian attains, is open.

A common budget for the same tensor. The two joint sixth-moment bounds share one tensor and scalar argument in `JointSixthCore.lean`: Hilbert–Schmidt and cubic-trace bounds, a completed square, and Young's inequality. The production cuts and parameterized refinements specialize that argument through the same real-to-complex normalization identities. The sixth moment and the Stein fluctuations cannot be optimized independently: both depend on the same contraction invariants. We now formalize this connection in Lean, including the real/complex normalization and a finite checker that reserves part of the contraction budget before accepting a Stein cut. Signed reserve coefficients are retained. Ordinary and joint Gram budgets specialize one signed-reserve soundness theorem. Likewise, ordinary spectral cuts specialize the general shifted-cut theorem, and Stein and spectral nodes use the same polynomial-majorant integration lemma for an arbitrary integrable cut with nonnegative expectation. The endpoint uses 41 such nodes, with each node restricted to its proved kurtosis range; the full checker and analytic bridge are in Lean. Retaining the upper bound on the negative cumulant contraction yields a stronger parameterized moment family. An exact transport identity then upgrades a node's majorant when its budget permits a larger moment-cut weight, without a new half-line search. These two refinements are also formalized in Lean. The moment costs also admit a convex formulation with two 2×2 PSD blocks and linear constraints. An exact rational checker connects each accepted recipe to the actual cubic moment inequality, replacing the outer parameter grid. Its exact scalar boundary is also proved in Lean. For moment weight $w > 0$ and s_0 cost $A > 0$, the minimum h_0 cost is

$$B_{\min}(w,A) = \begin{cases} 11781w/32 + 18225w^2/(4A), & 11A \le 360w, \\ 198w + 22500w^2/(A + 40w), & 11A \ge 360w. \end{cases}$$

Both branches have explicit feasible recipes. A three-rational checker for (w, A, B) therefore replaces the auxiliary moment-cost variables, and Lean connects its acceptance to the actual cubic on the checked node range. The contraction reserve and pointwise majorant still need their own exact checks. This is optimality within the certificate family; it does not identify all attainable response laws.

4.3 The remote functional

Let $L = z_1(a + b|z_1|^2)$ with the constants of the witness, and for a cubic G put $\Lambda(G) = \mathbb{E}|L+G| + \mathbb{E}|L-G| - \mathbb{E}|G|^2/\epsilon$ and $S = \sup_v\{2\mathbb{E}|L+\sqrt{v}U| - v/\epsilon\}$, U a complex Gaussian. The conjecture $\Lambda(G) \le S$ for all cubics ("Conjecture K") is the analytic heart of the lower bound.

- **A one-function reformulation [P].** Circular symmetry, $\mathbb{E}|W| = (\pi/2)\mathbb{E}|\Re W|$, and the identity $|a+b| + |a-b| = 2\max(|a|,|b|)$ for real a, b give exactly

$$\Lambda(G) = 2\mathbb{E}|L| + \pi\,\mathbb{E}\big(|\Re G| - |\Re L|\big)_+ - \mathbb{E}|G|^2/\epsilon,$$

  so $\sup_G \Lambda - 2\mathbb{E}|L| = \sup_t\big[\frac{\pi^2\epsilon}{4}\|P_{2,1}t\|^2 - \pi\mathbb{E}(|t|\,|\Re L|)\big]$ over real functions $|t| \le 1$. All of the gain sits near the hyperplane $\Re z_1 = 0$.

- **The Gaussian is a strict local maximum [P,C].** At the Gaussian point every direction that involves z_1 has negative second variation, certified in interval arithmetic for seven sector coefficients. The Hessian's kernel is the entire z_1-independent cubic sector; a negative fourth-order term controls those genuine cubic directions.

- **The independent sector [P,C].** For every cubic q not involving z_1, $\Lambda(q) < S + 3.6\cdot 10^{-7}$. Moment inequalities alone stall about 10^{-5} above S, because $h(x) = 2\mathbb{E}|L + \sqrt{x}|$ is a Bernstein function and the three-moment relaxation is maximised by a two-point law [P]. Going below that floor needs the Stein and spectral cuts.

- **Two natural shortcuts are false [P,C].** $|q|^2$ does not dominate $\mathbb{E}|q|^2 \cdot \mathrm{Exp}(1)$ in convex order, and there is no Gaussian Laplace domination: $Q = X(4(|X|^2-2)+3i(|Y|^2-1))/\sqrt{41}$ has $\mathbb{E}e^{-20|Q|^2} < 0.0453 < 1/21$.

4.4 Actual responses

Instead of bounding Λ for all cubics, one can use that the cubic comes from a bounded response. Write F for the cubic part of the response and $\mathcal{N}$ for the value of the witness on it.

- **Exact support [P].** The paired support $\frac{1}{2}(\mathbb{E}|A+B| + \mathbb{E}|A-B|)$ is exactly the pointwise support function of the complex response set $\{(h,k) = (e^{i\alpha}\cos\theta,\, ie^{i\alpha}\sin\theta)\}$, because $(|A+B| + |A-B|)^2 = 4\max_\theta |A\cos\theta - iB\sin\theta|^2$.

- **Concentrated cubics are harmless [P,C].** At $\epsilon = 1891/2500$ on the central rows $0.84 \le t \le 0.851$, $\mathcal{N} < 0.71325 < 1/1.402$ whenever the cubic part of the response concentrates enough energy on a structured subspace:
  - 80% on one diagonal Wick direction;
  - 90% on one mixed monomial $X^2\bar{Y}$, or on a whole slice $\{X^2\bar{A}\}$;
  - 93.5%, 96.2%, 97.6%, 99.2% on any cubic in $2, 3, 4, 6$ remote coordinates;
  - 95.5% to 99.95% on $H_{2,0}(A)H_{0,1}(B)$ with $\dim A \le 24$, via a sharp support theorem whose extremiser has law $R_m G/\sqrt{m(m+1)}$, $R_m \sim \Gamma(m,1)$.

  The remainder of the cubic may be arbitrary and fully coupled.

- **The obstruction to covering [P].** $Q_N = X_1(\sum_{j\le N}|X_j|^2 - N - 1)/\sqrt{N+1}$ keeps an influential direction while its projection onto every fixed-dimensional core tends to zero. So "small core or Gaussian" is not an exhaustive dichotomy.

- **How far from the benchmark structured cubics sit [F].** Remote-only payoffs relative to S:

sum of N slices, $N=$	1	2	4	8	16	24	∞ (Gaussian)
$S - \sup\Lambda$ (10^{-3})	1.68	1.42	1.08	0.73	0.43	0.32	0

The tested concentrated cubics leave room of order 10^{-3}; these diffuse sums approach the benchmark. This evidence does not classify all near-maximizers.

4.5 Limits of the witness method

- **Ceilings [P,C] and [Lean].** Actual finite Gaussian responses give the scalar cubic family a ratio ceiling < 1.4022581233 when taking suprema over Gaussian dimensions, for every complex multiplier; its generic cancellation arithmetic is in Lean, its Gaussian construction and integral bounds are external. At $\epsilon = 149/200$ a genuine response gives norm > 0.71329, excluding 1.402 for that witness. Explicit "chirp" responses prove that games of degree ≤ 5 cannot exceed 1.4036705, and of degree ≤ 7 cannot exceed 1.4044895 [P,C].
- **Extending the pattern fails [P,C].** The full alternating game $\sum_{k\geq 0}(-1)^k P_{k+1,k}$ has $\sup_n \|\cdot\|_{\infty\to 1} = 1$, attained by explicit unimodular responses (one in complex dimension 308,835 is replayed). So better high-degree games must be re-optimised against new responses, not extrapolated.
- **Certified finite exchange [P].** For a finite response pool truncated at level D with value v_D, a global witness norm cap $v_D + \delta$ and a realizable mixture with omitted coefficient tail at most $\tau < v_D$ give $1/(v_D+\delta) \leq K_G^{\mathbb{C}} \leq 1/(v_D - \tau)$ by the two one-sided transfers. The reciprocal interval has width $\delta + \tau$; both the uniform witness cap and the complete mixture tail must actually be certified.
- **A global spherical form, and why it does not certify cheaply [P] and [Lean].** Finite Gaussian Hermite games transfer, in every Gaussian dimension, to compact operators on the sphere with the same bounded responses and the same shared drivers; the numerical bounds are unchanged. The dual operator can then be truncated with a rigorous error: the cubic multiplier tail and the signed cap assembly are proved in Lean (`SphericalTransfer.cubic_tail_32_12`, `global_cap_from_truncation`); the analytic bridge is a written proof (`docs/SPHERICAL_TRANSFER.md`). But the operator has no square-integrable kernel: its multiplicity-weighted squared coefficients grow like r^7. So no Hilbert–Schmidt-type certificate can control it, and direct coefficient optimisation is impractical (about $8 \cdot 10^{20}$ coefficients in the tested truncation). A useful norm cap must bound the signed action on bounded responses, exploiting the unitary symmetry.

5 Why the gap is hard

The witness class caps the lower bound. A game supported on degrees $\{1, 3\}$ sees only the cubic constraint on roundings. Even a perfect evaluation of the norm of $P_{1,0} - \epsilon P_{2,1}$ cannot certify more than 1.4022581233 (the certified family ceiling above). Proved obstructions are known for higher degrees: every game of degree ≤ 5 is capped at 1.4036705, and every game of degree ≤ 7 at 1.4044895, both below Haagerup's value. Reaching the upper end of the interval therefore needs games of degree ≥ 9; in the real case the certified game uses all odd degrees up to 21.

Certifying higher degree is the real bottleneck. Two different certificate architectures, gated rows and support functions, both lose about $2 \cdot 10^{-2}$ on the best degree-5 game. The reason is structural: a bounded function with a large cubic part is forced to have a degree-five part too. Through the visible coordinate, the cubic part sees the average of the remote coefficient profile

and the quintic part its variation, and boundedness makes the profile vary. A certificate that treats the chaos degrees as independent cannot see this coupling. The coupling is *driver sharing*: in the dimension-limit model the best degree-5 game has norm 0.712464 ($1/N = 1.40358$) when its degree-5 mixed sectors $(|z_1|^2 - 1)W$ and $z_1^2\overline{W}$ are driven by the same remote cubic W as the cubic sector, and 0.724753 when they are driven by independent copies; the optimum then switches the degree-5 part off. So a higher-degree certificate must let one remote residual feed several degrees through the visible coordinate. The cubic rows already do this for their degree-3 sectors; extending them to the degree-5 companions of the same residual is the next construction [F].

Why the complex case is harder than the real one. For real ± 1 responses, $h = \frac{1}{2}(f + g)$ and $k = \frac{1}{2}(f - g)$ have disjoint supports, so the two sides of a game decouple at every degree. For complex unimodular responses $|h|^2 + |k|^2 = 1$ and $\Re(h\bar{k}) = 0$ force $f = e^{i\alpha}e^{i\theta}$, $g = e^{i\alpha}e^{-i\theta}$: one shared phase field α drives both sides, and they never separate. The same shared phase is why marginal laws do not determine the relevant nonlinear correlations: $(H(X), X)$ and $(H(Y), X)$, with $H(X) = X(|X|^2 - 2)/\sqrt{2}$ and X, Y independent, have identical marginals and covariances, but $\mathbb{E}[\bar{r}\, q|q|^2]$ equals $8\sqrt{2}$ in the first case and 0 in the second.

6 Open problems

Planned work is tracked as GitHub issues of the repository: epics #12 (unconditional lower bound), #17 (upper inputs), #22 (better lower bound), #29 (build infrastructure) and #34 (research plans).

1. Complete and verify the node caps (`CapsRest`, the remaining vertex-cap hypothesis), which would make the lower bound fully machine-checked. The caps of rows 12–15 are localised caps: coverage by constrained pieces is checked in Lean, and most pieces pass; the remaining pieces fail the exact positivity check of their S-lemma matrix by about 10^{-5} and need re-optimised multipliers.
2. Discharge the sharp upper rounding's numerical comparison and energy-table error; its Hermite expansion, structural analytic bridge and Hardy tail bound are formalized.
3. Certify a game of degree ≥ 5 with a *multi-grade fibre*: one bounded function of the visible coordinate and one remote direction per degree, so that the cross-degree coupling is kept. This is the route to a lower bound above 1.403.
4. Run a complex primal/dual search, alternating linear programs over games of degree $\leq D$ and pools of roundings, to locate $K_G^{\mathbb{C}}$ to about 10^{-4} and to test whether Haagerup's value $1/\gamma^\dagger$ is the answer.
5. Prove the sharp first-moment bound $\mathbb{E}|q| \leq \frac{\sqrt{\pi}}{2}\|q\|_2$ for cubic Wick polynomials. It is natural, it appears to be open, and it is the largest single lever on the law side.

How this was made

The constructions and certificates were found by several AI research systems working on a shared manuscript, cross-checking each other's packages by independent replay. The endpoint theorems have the explicit Lean hypothesis lists in Section 3, and their axioms are audited in CI. The discoveries labelled [P] and [C] also include paper proofs and external exact or interval certificates; they are not all formalized in Lean. Selected written proofs and reproducible certificate records accompany the repository; the larger working manuscript is maintained separately.

References

- A. Grothendieck, *Résumé de la théorie métrique des produits tensoriels topologiques*, 1953.
- U. Haagerup, *A new upper bound for the complex Grothendieck constant*, Israel J. Math. 60 (1987).
- A. M. Davie, *Lower bound for K_G*, unpublished note, 1984.
- S. Heilman, C. Jones, G. Malavolta, *Sharper bounds for the complex Grothendieck constant*, arXiv:2609.25462.
- S. Heilman, C. Jones, G. Malavolta, *Primal/dual method for the Grothendieck constant*, arXiv:2610.10477.
- Guo, Fang, Lu, arXiv:2609.07000.
- R. Saha et al., *New lower and upper bounds for the Grothendieck constant*, arXiv:2608.11158.
- A. Naor, O. Regev, *Krivine schemes are optimal*, Proc. AMS 142 (2014), arXiv:1205.6415.
- G. Pisier, *Grothendieck's theorem, past and present*, Bull. AMS 49 (2012), arXiv:1101.4195.
- N. Brunner et al., *Bell nonlocality*, Rev. Mod. Phys. 86 (2014), arXiv:1303.2849.

Packed source certificate. A strict structural token reader agrees with the unchanged public dyadic parser, including malformed input. The generic delimiter identity covers legacy single-character splitting. Bounded character-list certificates assemble all 460 records of the last first-center source chunk and identify its public decoded result with exact integer witnesses. Source equality is tied to the original packed array, and the complete certificate is checked in the default root. Whole-array source identities, replacement of native band checks and Gaussian approximation estimates remain separate obligations.

Phase source identity. The identity expressing the squared production phase in Gram coordinates is available independently of the numerical phase cover. This lets local source certificates use the actual rounding polynomial before cover assembly.

Production source error propagation. A supplied integer leaf check bounds the production phase on that leaf; no concrete leaf check is supplied here. The existing native cover already bounds the phase globally. A shared Gram-coordinate identity supplies the domain constraint for both routes. On $[-1, 1]$ the cubic phase factor is $5/2$-Lipschitz. For positive radial parameter, the rounding-slice error is bounded by the quadratic Gaussian-base remainder plus the phase approximation error. Its phase bound remains an explicit premise to preserve independence from the native cover. Numerical source comparisons and integral certificates still have to instantiate these analytic bounds.

The conditional Gaussian L^2 triangle inequality on one unsmoothed radial slice charges the base remainder once, giving $5\rho^2$ plus $5/2$ times the phase-model error and the model-to-comparison error. The actual phase bound follows from the existing native cover; it remains a premise to keep these analytic theorems independent of that native computation. Choosing the actual phase as the bounded comparison phase makes its phase error zero. The remaining Gaussian model-to-comparison integral and the comparison function's relation to stored coefficients still need certificates. Radial integration, smoothing and band projection are separate steps toward the endpoint.

Keeping the actual phase inside the Gaussian model removes the phase-error term exactly. The remaining saved-model discrepancy is one normalized conditional Gaussian integral of model energy, saved energy and their cross term. Separate energy upper bounds U, V and a real cross lower bound L give the squared-discrepancy budget $U + V - 2L$. A complex cross enclosure with center c and radius r gives $U + V - 2\,\mathrm{Re}(c) + 2r$. These numerical receipts, the saved-source binding and outer radial weights remain separate obligations.

Rational source evaluation. The literal rational source polynomial is identified with the actual polynomial under diagonal gate scaling, including its coefficients and support. A finite moment-scaling identity preserves both Gaussian Jacobian factors, and a rational evaluator casts to the existing correlated-moment formula. The actual gate's outer moments specialize to rational Schur coefficients with inverse mass $1/(RQ)$ and positivity from its checked gate. Conditional-polynomial scaling and the eight source

coefficient factors combine into a full rational Gram sum, retaining conjugation, every cross term, and one kernel mass and Jacobian factor per physical pair. Complete source proofs identify the actual coherent norm with this value and bound the infinite weighted matrix energy by it. Numerical source bounds and the final endpoint remain open; no archived estimate is accepted as a proof.

Finite charge decomposition. The rational moment balance test makes unequal signed source charges orthogonal in the actual conditional Gram evaluator. Thus its full rational source norm equals the sum with cross-charge entries removed. Polynomial insertion is a rational linear map, so the actual source can also be split using the finite charge set of the original multiplier before expanding its eight source variables. The same insertion reverses physical charge, including when its rational rate is zero. Bilinearity and charge orthogonality therefore give the exact identity

$$N_g = \sum_{k \in \mathrm{chargeSet}(P_g)} B_g\big(I_{r_g}(P_{g,k}), I_{r_g}(P_{g,k})\big),$$

where N_g is the actual rational source norm expression, $P_{g,k}$ its physical multiplier's charge component, I_{r_g} the existing coherent insertion, and B_g the existing rational tensor pair kernel with both source coefficient factors. No charge count or square formula is assumed here. Laguerre coefficient binding, block diagonalization into squares, and numerical source enclosures remain separate obligations.

Zero and degree-one covariance contractions. The actual rational gate evaluator has exact zero and degree-one contraction formulas under its original checked-gate premise. Its zero contraction is the pair mass n^2/δ; the outer and inner degree-one squared contractions are this mass times χ_+ and χ_-. The ground-versus-quadratic contraction is the mass times λ, retaining the intrinsic inserted constant n, the conditional variance $R + P$, and the Jacobian $1/(RQ)$. This is not an equality between opposite-charge Hermitian Gram entries. Higher-degree diagonalization, the complete 18-block/1140-square identity, numerical source and energy enclosures, and the final endpoint remain open.